

NAME: Triveni Kallappa Biradar

PRN : 202401070172

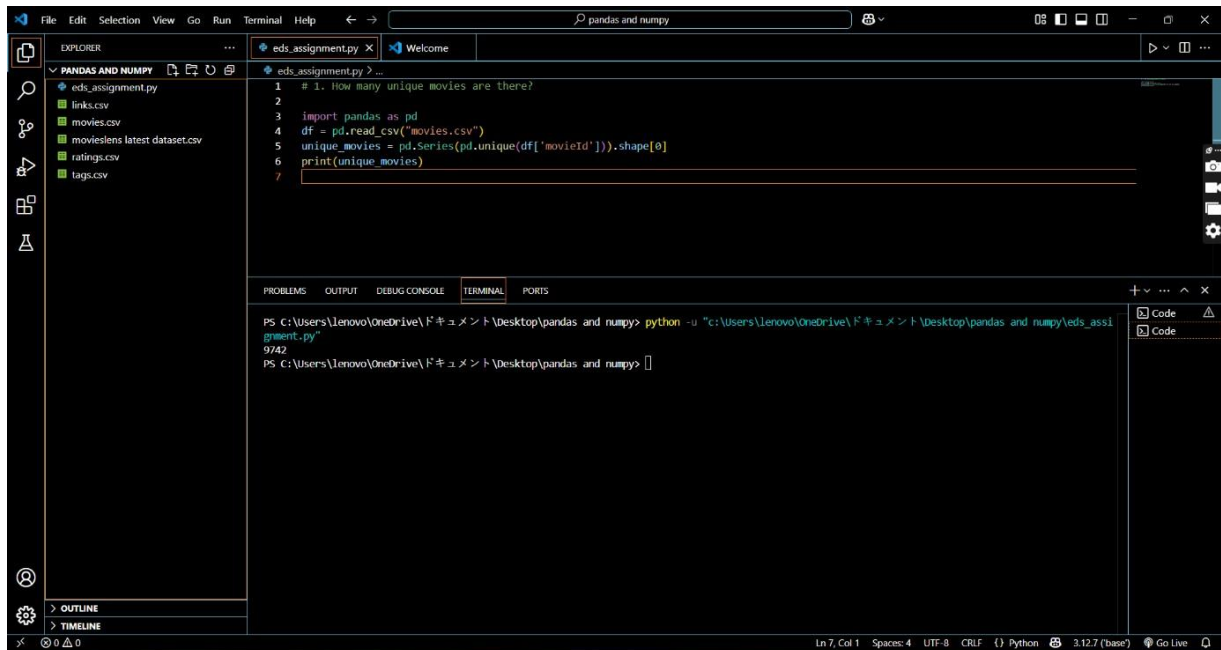
ROLLNO : ET2-44

CLASS : ET2

EDS THEORY ASSIGNMENT 1

URL: <https://www.kaggle.com/datasets/shubhammehta21/movie-lens-small-latest-dataset>

1. How many unique movies are there?



The screenshot shows the Visual Studio Code interface with a file explorer on the left containing 'PANDAS AND NUMPY' files: 'eds_assignment.py', 'links.csv', 'movies.csv', 'movielens latest dataset.csv', 'ratings.csv', and 'tags.csv'. The main editor displays 'eds_assignment.py' with the following code:

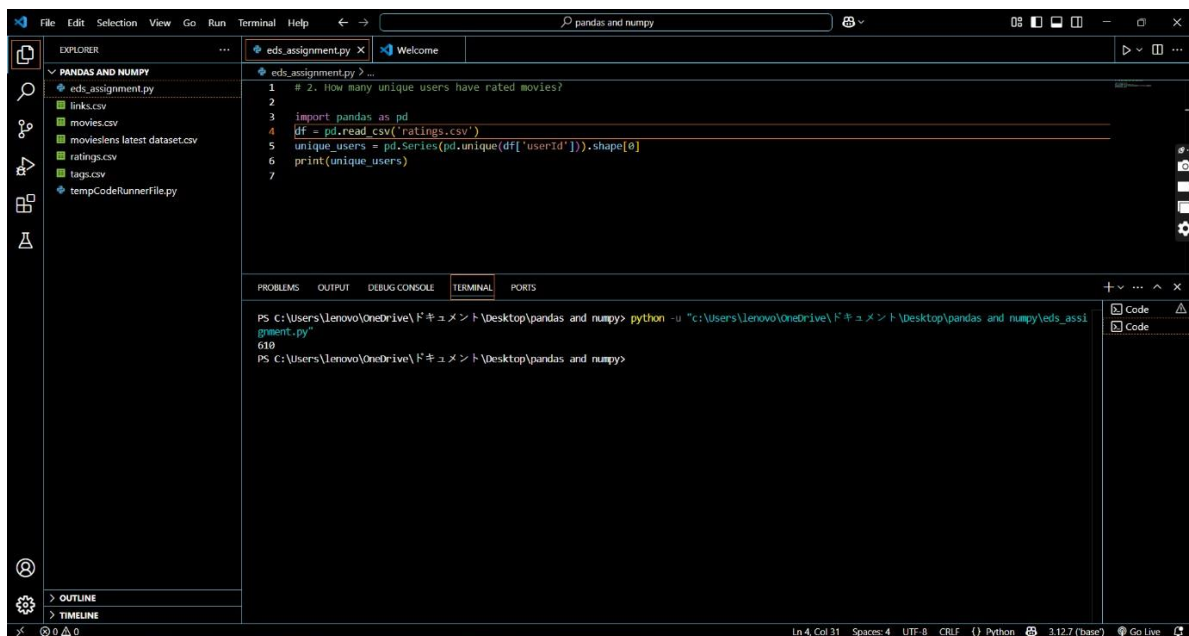
```
1 # 1. How many unique movies are there?
2
3 import pandas as pd
4 df = pd.read_csv("movies.csv")
5 unique_movies = pd.Series(pd.unique(df['movieId'])).shape[0]
6 print(unique_movies)
7
```

The bottom panel shows the TERMINAL output:

```
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"
9742
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```

The status bar at the bottom indicates 'Ln 7, Col 1', 'Spaces: 4', 'UTF-8', 'CRLF', 'Python', and '3.12.7 (base)'.

2. How many unique users have rated movies?



The screenshot shows the Visual Studio Code interface with a file explorer on the left containing 'PANDAS AND NUMPY' files: 'eds_assignment.py', 'links.csv', 'movies.csv', 'movielens latest dataset.csv', 'ratings.csv', 'tags.csv', and 'tempCodeRunnerFile.py'. The main editor displays 'eds_assignment.py' with the following code:

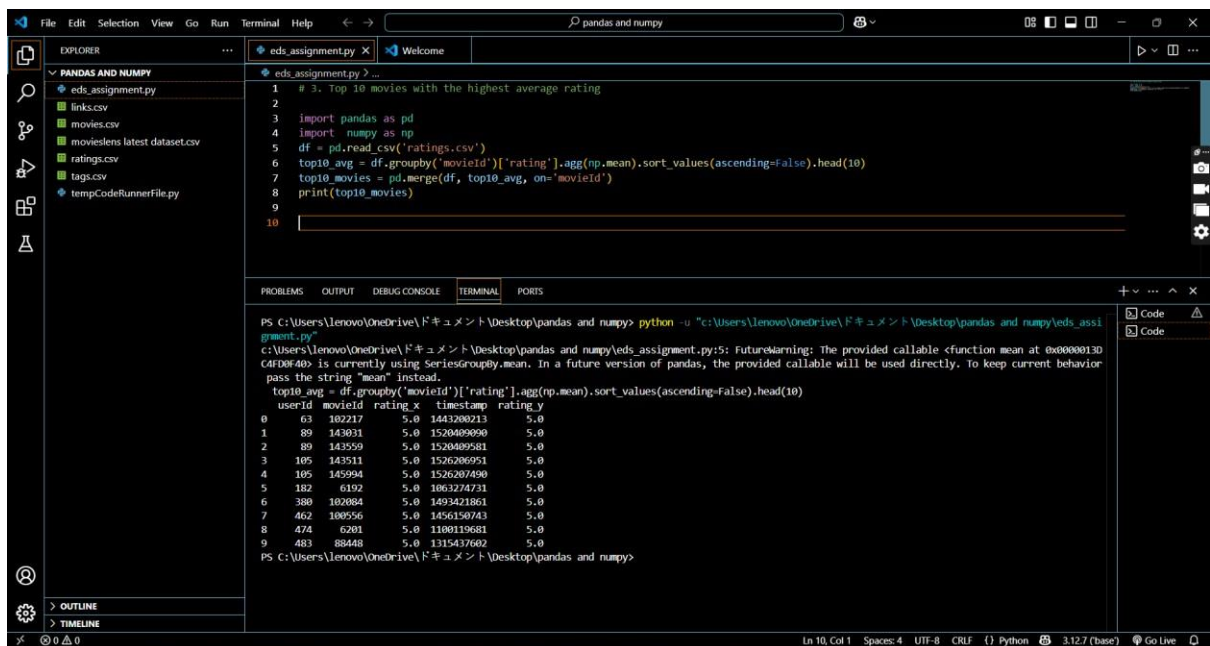
```
1 # 2. How many unique users have rated movies?
2
3 import pandas as pd
4 df = pd.read_csv('ratings.csv')
5 unique_users = pd.Series(pd.unique(df['userId'])).shape[0]
6 print(unique_users)
7
```

The bottom panel shows the TERMINAL output:

```
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"
610
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```

The status bar at the bottom indicates 'Ln 4, Col 31', 'Spaces: 4', 'UTF-8', 'CRLF', 'Python', and '3.12.7 (base)'.

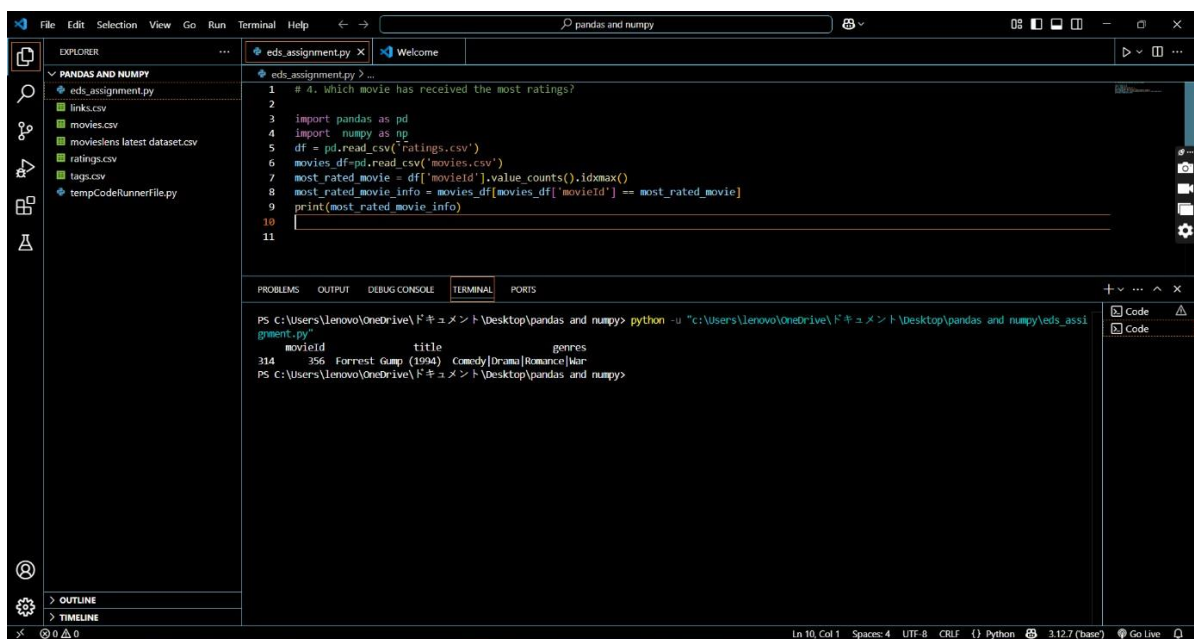
3. Top 10 movies with the highest average rating



```
1 # 3. Top 10 movies with the highest average rating
2
3 import pandas as pd
4 import numpy as np
5 df = pd.read_csv('ratings.csv')
6 top10_avg = df.groupby('movieId')['rating'].agg(np.mean).sort_values(ascending=False).head(10)
7 top10_movies = pd.merge(df, top10_avg, on='movieId')
8 print(top10_movies)
9
10
```

```
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"
c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py:5: FutureWarning: The provided callable <function mean at 0x00000130C4FD0840> is currently using SeriesGroupBy.mean. In a future version of pandas, the provided callable will be used directly. To keep current behavior pass the string "mean" instead.
top10_avg = df.groupby('movieId')['rating'].agg(np.mean).sort_values(ascending=False).head(10)
  userId  movieId  rating_x  timestamp  rating_y
0      63   102217      5.0   1443200213      5.0
1      89   143031      5.0   1520409090      5.0
2      89   143559      5.0   1520409581      5.0
3     105   143511      5.0   1526206951      5.0
4     105   145994      5.0   1526207490      5.0
5     182    6192      5.0   1063274731      5.0
6     380   102084      5.0   1493421861      5.0
7     462   100556      5.0   1456150743      5.0
8     474    6201      5.0   1100119681      5.0
9     483   88448      5.0   1315437602      5.0
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```

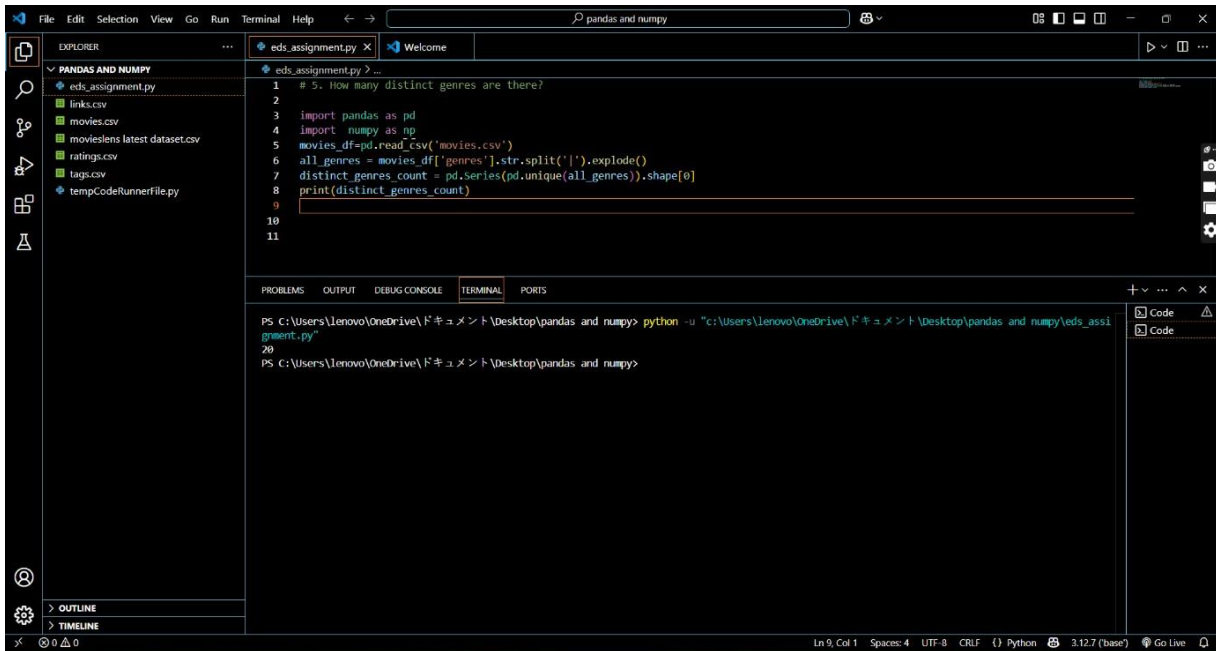
4. Which movie has received the most ratings?



```
1 # 4. Which movie has received the most ratings?
2
3 import pandas as pd
4 import numpy as np
5 df = pd.read_csv('ratings.csv')
6 movies_df = pd.read_csv('movies.csv')
7 most Rated movie = df['movieId'].value_counts().idxmax()
8 most Rated movie info = movies_df[movies_df['movieId'] == most Rated movie]
9 print(most Rated movie info)
10
11
```

```
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"
  movieId      title  genres
314    356  Forrest Gump (1994)  Comedy|Drama|Romance|War
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```

5. How many distinct genres are there?

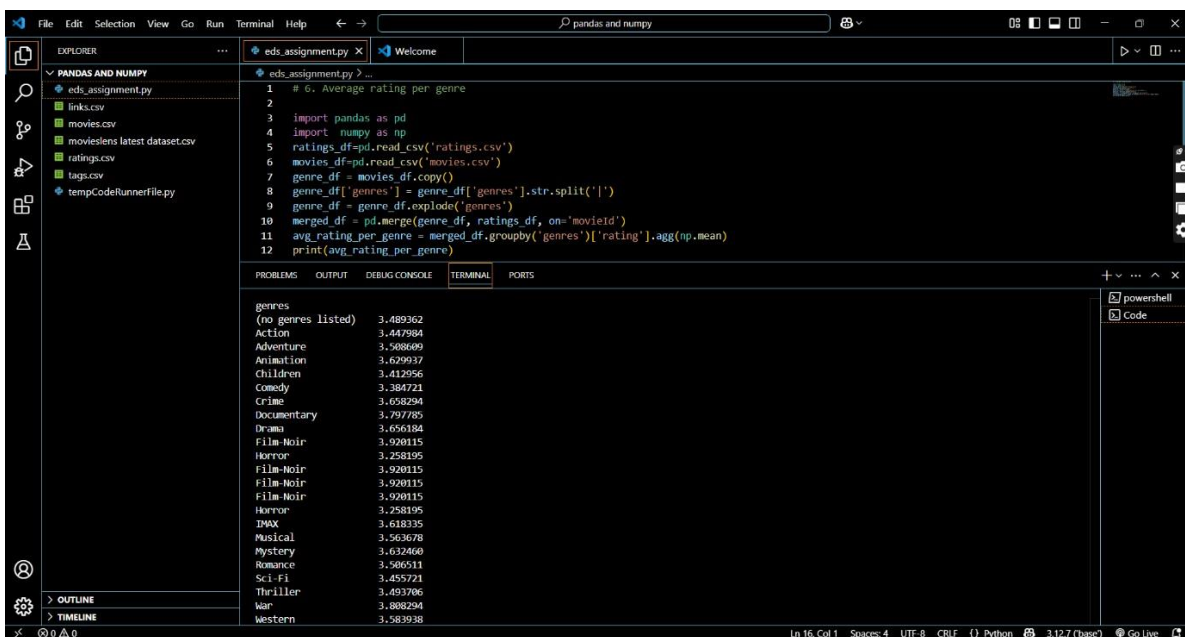


The screenshot shows a VS Code window with a file explorer on the left containing files like `links.csv`, `movies.csv`, `movielens latest dataset.csv`, `ratings.csv`, `tags.csv`, and `tempCodeRunnerFile.py`. The main editor displays `eds_assignment.py` with the following code:

```
1 # 5. How many distinct genres are there?
2
3 import pandas as pd
4 import numpy as np
5 movies_df = pd.read_csv('movies.csv')
6 all_genres = movies_df['genres'].str.split('|').explode()
7 distinct_genres_count = pd.Series(pd.unique(all_genres)).shape[0]
8 print(distinct_genres_count)
9
10
11
```

The terminal at the bottom shows the command `python -u "c:\Users\lenovo\OneDrive\Desktop\pandas and numpy\eds_assignment.py"` being executed, resulting in the output: `PS C:\Users\lenovo\OneDrive\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\Desktop\pandas and numpy\eds_assignment.py"` and `PS C:\Users\lenovo\OneDrive\Desktop\pandas and numpy>`.

6. Average rating per genre



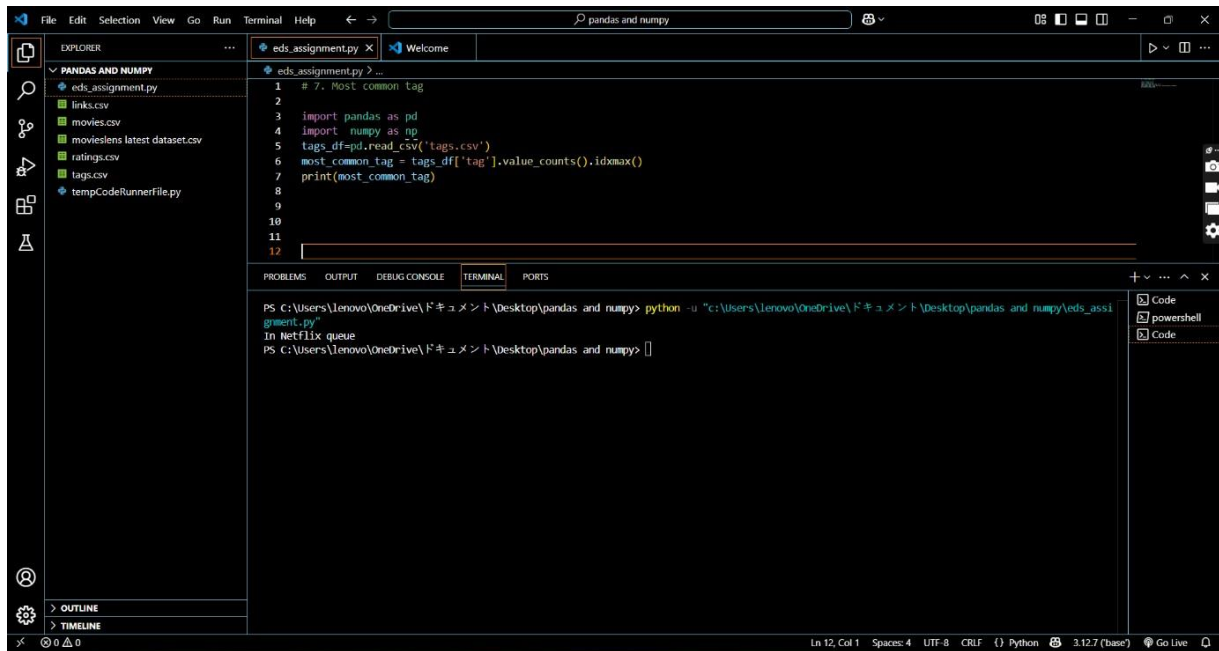
The screenshot shows a VS Code window with the same file explorer as before. The main editor displays `eds_assignment.py` with the following code:

```
1 # 6. Average rating per genre
2
3 import pandas as pd
4 import numpy as np
5 ratings_df = pd.read_csv('ratings.csv')
6 movies_df = pd.read_csv('movies.csv')
7 genre_df = movies_df.copy()
8 genre_df['genres'] = genre_df['genres'].str.split('|')
9 genre_df = genre_df.explode('genres')
10 merged_df = pd.merge(genre_df, ratings_df, on='movieid')
11 avg_rating_per_genre = merged_df.groupby('genres')['rating'].agg(np.mean)
12 print(avg_rating_per_genre)
```

The terminal at the bottom shows the output of the code, which is a table of average ratings per genre:

genres	(no genres listed)
Action	3.489362
Adventure	3.447984
Animation	3.508689
Children	3.629937
Comedy	3.412956
Crime	3.384721
Documentary	3.658294
Drama	3.797785
Film-Noir	3.656184
Horror	3.920115
Horror	3.258195
Horror	3.920115
Horror	3.920115
Horror	3.920115
Horror	3.258195
IMAX	3.618335
Mystery	3.563678
Mystery	3.632460
Romance	3.506511
Sci-Fi	3.455721
Thriller	3.493706
War	3.808294
Western	3.583938

7. Most common tag



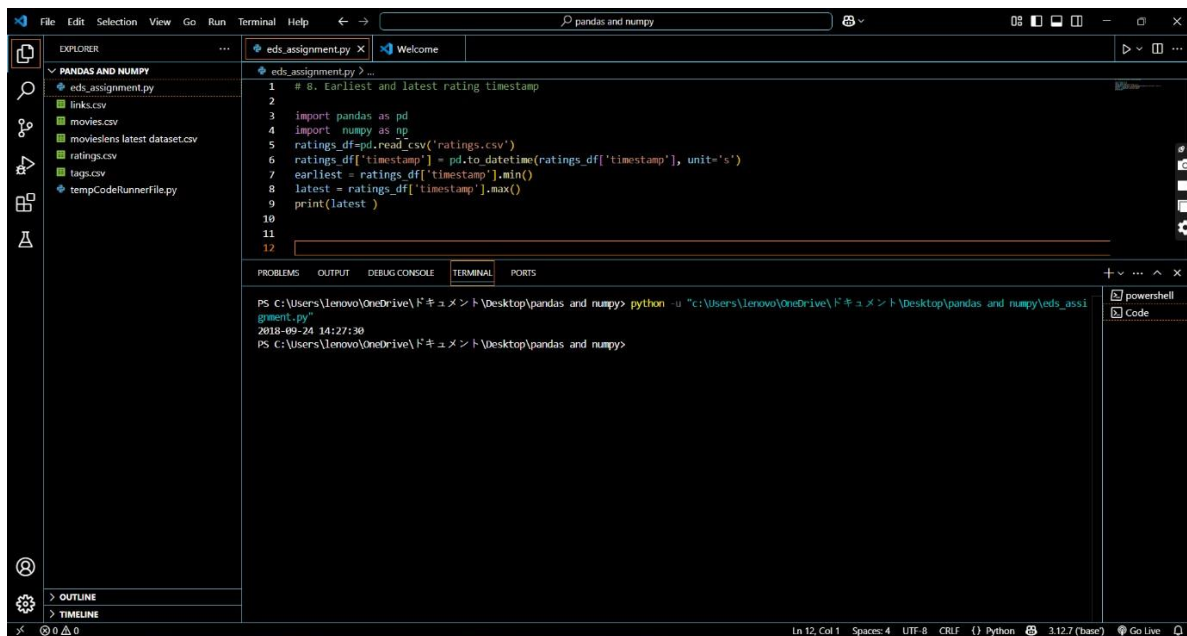
The screenshot shows the Visual Studio Code editor with a file explorer on the left containing a folder named 'PANDAS AND NUMPY'. Inside this folder are files: 'eds_assignment.py', 'links.csv', 'movies.csv', 'movielens latest dataset.csv', 'ratings.csv', 'tags.csv', and 'tempCodeRunnerFile.py'. The 'eds_assignment.py' file is open in the editor, showing the following Python code:

```
1 # 7. Most common tag
2
3 import pandas as pd
4 import numpy as np
5 tags_df=pd.read_csv('tags.csv')
6 most_common_tag = tags_df['tag'].value_counts().idxmax()
7 print(most_common_tag)
8
9
10
11
12
```

The bottom panel of the editor shows the 'TERMINAL' tab with the following output:

```
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"
In Netflix queue
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```

8. Earliest and latest rating timestamp



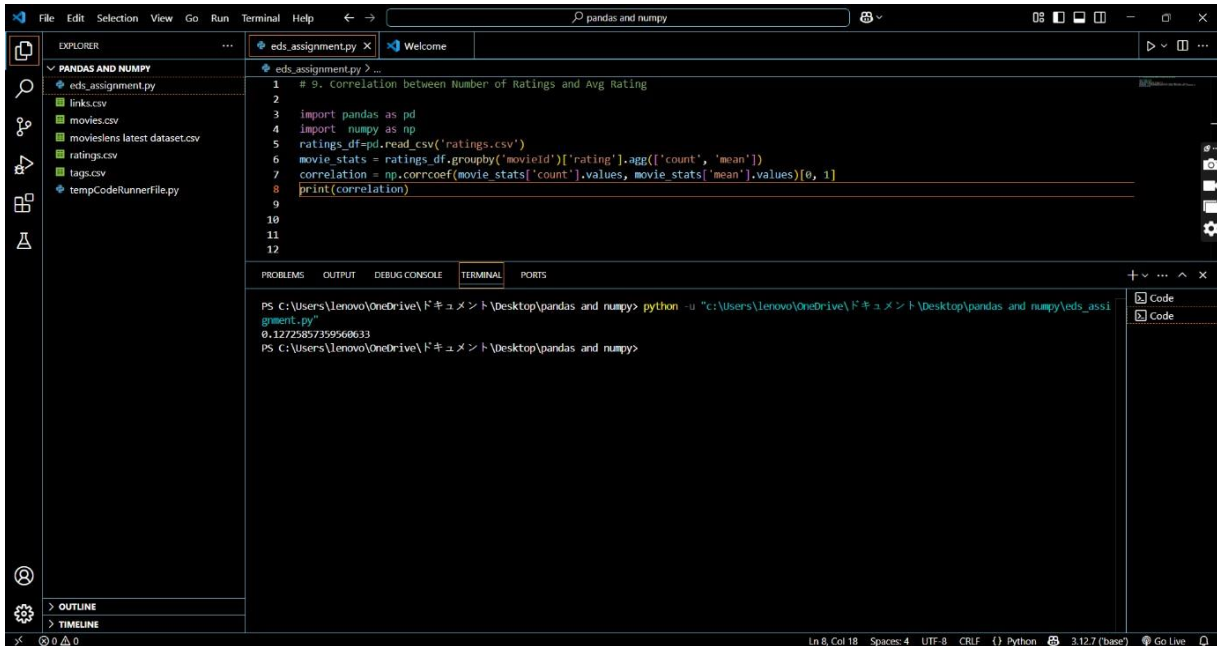
The screenshot shows the Visual Studio Code editor with the same file explorer as in the previous image. The 'eds_assignment.py' file is open, showing the following Python code:

```
1 # 8. Earliest and latest rating timestamp
2
3 import pandas as pd
4 import numpy as np
5 ratings_df=pd.read_csv('ratings.csv')
6 ratings_df['timestamp'] = pd.to_datetime(ratings_df['timestamp'], unit='s')
7 earliest = ratings_df['timestamp'].min()
8 latest = ratings_df['timestamp'].max()
9 print(earliest)
10
11
12
```

The bottom panel of the editor shows the 'TERMINAL' tab with the following output:

```
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"
2018-09-24 14:27:30
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```

9. Correlation between number of ratings and average rating



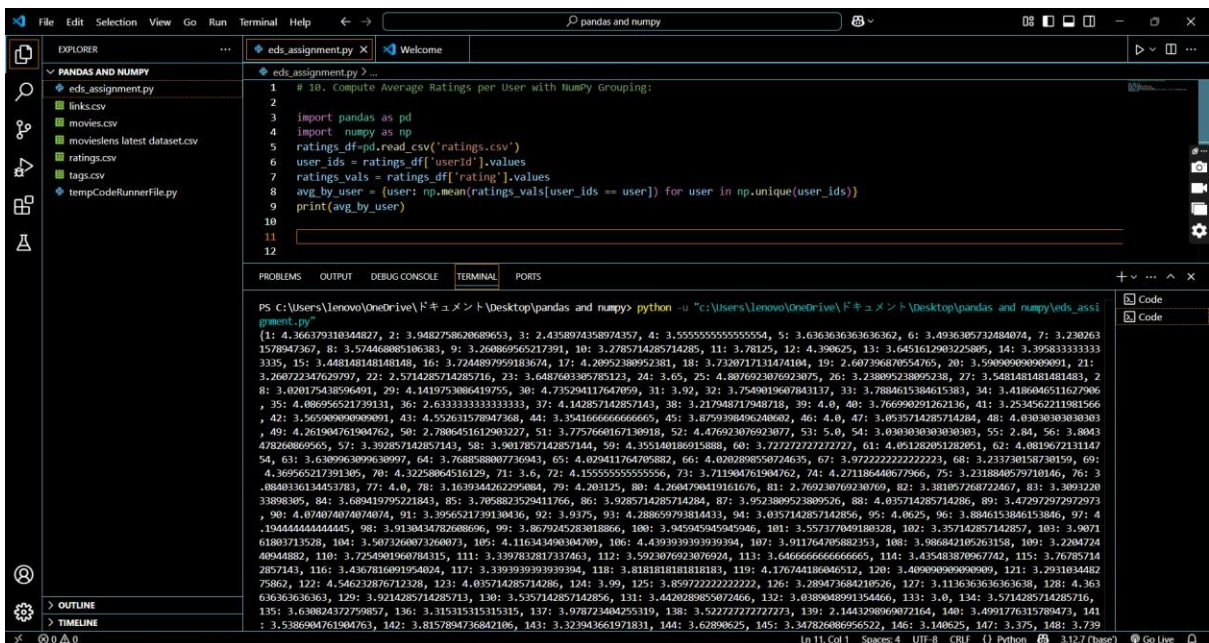
```
1 # 9. Correlation between Number of Ratings and Avg Rating
2
3 import pandas as pd
4 import numpy as np
5 ratings_df = pd.read_csv('ratings.csv')
6 movie_stats = ratings_df.groupby('movieId')['rating'].agg(['count', 'mean'])
7 correlation = np.corrcoef(movie_stats['count'].values, movie_stats['mean'].values)[0, 1]
8 print(correlation)
9
10
11
12
```

PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"

0.1272585735956633

PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>

10. Compute Average Ratings per User with NumPy Grouping:



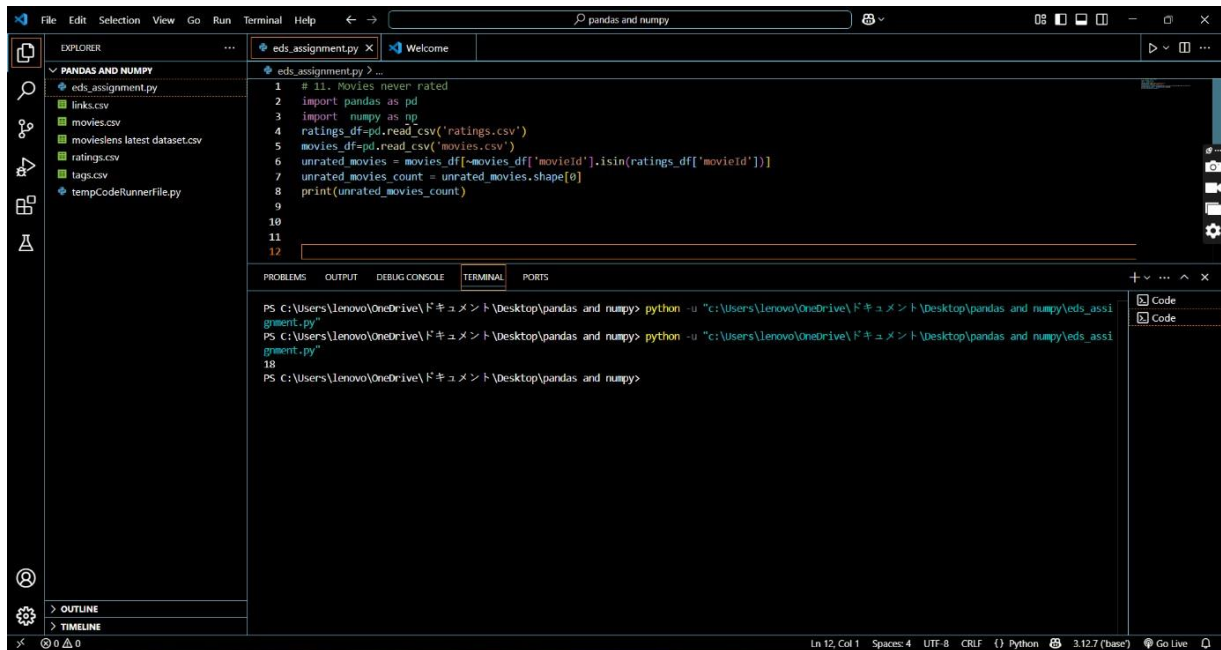
```
1 # 10. Compute Average Ratings per User with NumPy Grouping:
2
3 import pandas as pd
4 import numpy as np
5 ratings_df = pd.read_csv('ratings.csv')
6 user_ids = ratings_df['userId'].values
7 ratings_vals = ratings_df['rating'].values
8 avg_by_user = {user: np.mean(ratings_vals[user_ids == user]) for user in np.unique(user_ids)}
9 print(avg_by_user)
10
11
12
```

PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"

{1: 4.366379318344827, 2: 3.948275862689653, 3: 2.4358974358974357, 4: 3.5555555555555554, 5: 3.6363636363636362, 6: 3.4936305732484074, 7: 3.2302631578947367, 8: 3.574468085186383, 9: 3.26089565217391, 10: 3.2785714285714285, 11: 3.78125, 12: 4.390625, 13: 3.6451612903225805, 14: 3.395833333333333, 15: 3.448148148148148, 16: 3.7244897959183674, 17: 4.2052388952381, 18: 3.7280717131474104, 19: 2.667366070554705, 20: 3.5089898989898991, 21: 3.26072247628797, 22: 2.5714285714285716, 23: 3.6487692307587122, 24: 3.65, 25: 4.8276923076923075, 26: 3.238095238095238, 27: 3.5481481481481483, 28: 3.020175438596491, 29: 4.1419753086419755, 30: 4.735204117647059, 31: 3.92, 32: 3.7549019607843137, 33: 3.7884615384615383, 34: 3.418606511627906, 35: 4.086956521739131, 36: 2.6333333333333333, 37: 4.142857142857143, 38: 3.217948717948718, 39: 4.0, 40: 3.766990291262136, 41: 3.2534562211881566, 42: 3.565909090909091, 43: 4.552631578947368, 44: 3.3541666666666665, 45: 3.8759398496240602, 46: 4.0, 47: 3.0535714285714284, 48: 4.0303030303030303, 49: 4.261904761904762, 50: 2.7806451612903227, 51: 3.7757660167130918, 52: 4.476923076923077, 53: 5.0, 54: 3.0303030303030303, 55: 2.84, 56: 3.8043478260869565, 57: 3.392857142857143, 58: 3.9017857142857144, 59: 4.355140186915888, 60: 3.7272727272727272, 61: 4.051282051282051, 62: 4.081967213114754, 63: 3.6309630963096309, 64: 3.768858807736943, 65: 4.05411764705882, 66: 4.0200898507744635, 67: 3.9722222222222223, 68: 3.233781587360159, 69: 4.3696652173913065, 70: 4.3258964516129, 71: 3.6, 72: 4.1555555555555556, 73: 3.711904761904762, 74: 4.271186408577966, 75: 3.231884857910146, 76: 3.084013613453783, 77: 4.0, 78: 3.163934426209084, 79: 4.281125, 80: 4.2694790419161676, 81: 2.769230769230769, 82: 3.381957268722467, 83: 3.300322033898305, 84: 3.689419795221843, 85: 3.705882329411766, 86: 3.9285714285714284, 87: 3.9523809523809526, 88: 4.035714285714286, 89: 3.4729727272727273, 90: 4.074074074074074, 91: 3.3956521739130436, 92: 3.9375, 93: 4.288659793814413, 94: 3.0357142857142856, 95: 4.0625, 96: 3.8846153846153846, 97: 4.1944444444444445, 98: 3.9130434782608696, 99: 3.8679245283018866, 100: 3.945945945945946, 101: 3.557377049180328, 102: 3.357142857142857, 103: 3.907161801713528, 104: 3.5073260873260873, 105: 4.116343490304709, 106: 4.4393939393939394, 107: 3.911764705882353, 108: 3.986842105263158, 109: 3.2204724094482, 110: 3.7254901960784315, 111: 3.339781281737863, 112: 3.5923076923076924, 113: 3.6466666666666665, 114: 3.435483870967742, 115: 3.767857142857143, 116: 3.4367816901954024, 117: 3.3939393939393939, 118: 3.8181818181818183, 119: 4.176784118606512, 120: 3.409090909090909, 121: 3.293103448275862, 122: 4.546232876712328, 123: 4.035714285714286, 124: 3.99, 125: 3.8592722222222222, 126: 3.289473684210526, 127: 3.1136363636363638, 128: 4.363636363636363, 129: 3.9214285714285713, 130: 3.5357142857142856, 131: 3.4420289859072466, 132: 3.0389048991354466, 133: 3.0, 134: 3.5714285714285716, 135: 3.608024372759857, 136: 3.315315315315315, 137: 3.978723404255319, 138: 3.522727272727273, 139: 2.1443298969072164, 140: 3.4991776315789473, 141: 3.5386904761904763, 142: 3.8157894736842106, 143: 3.323943661971831, 144: 3.62890625, 145: 3.347826086956522, 146: 3.140625, 147: 3.375, 148: 3.739}

Ln 11, Col 1 Spaces: 4 UTF-8 CRLF () Python 3.12.7 (base) Go Live

11. Movies never rated



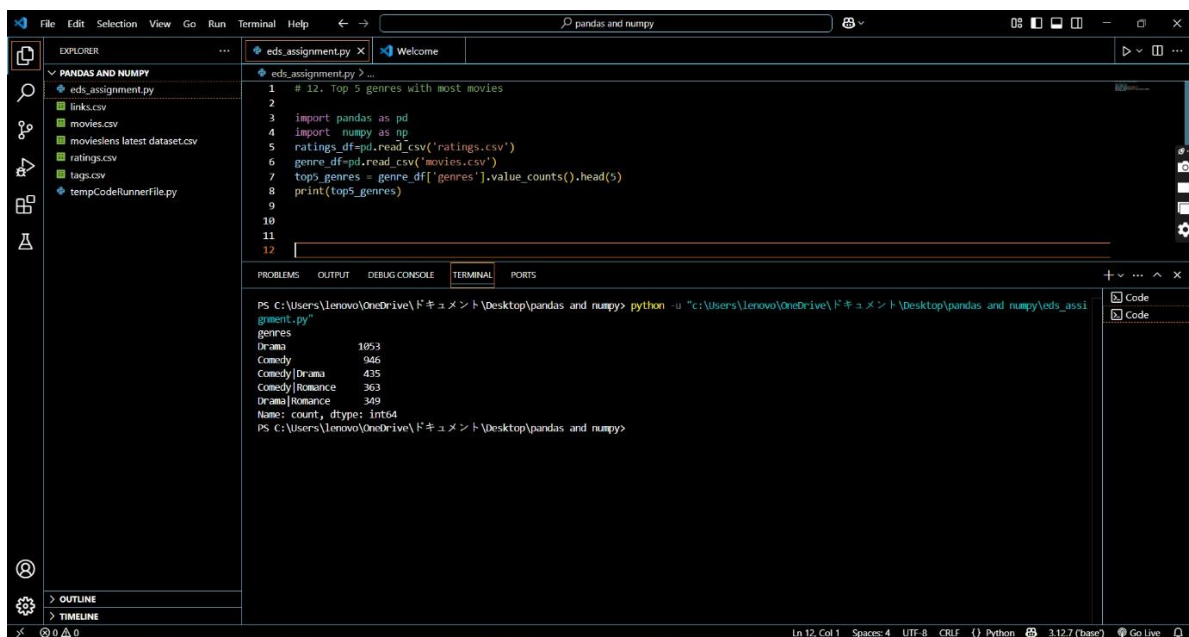
The screenshot shows a VS Code editor window with a file explorer on the left and a code editor in the center. The file explorer shows a project named 'PANDAS AND NUMPY' with files: 'eds_assignment.py', 'links.csv', 'movies.csv', 'movielens latest dataset.csv', 'ratings.csv', 'tags.csv', and 'tempCodeRunnerFile.py'. The code editor shows the following Python code:

```
1 # 11. Movies never rated
2 import pandas as pd
3 import numpy as np
4 ratings_df = pd.read_csv('ratings.csv')
5 movies_df = pd.read_csv('movies.csv')
6 unrated_movies = movies_df[~movies_df['movieid'].isin(ratings_df['movieid'])]
7 unrated_movies_count = unrated_movies.shape[0]
8 print(unrated_movies_count)
9
10
11
12
```

The terminal output shows the command to run the script and the result:

```
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"
18
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```

12. Top 5 genres with most movies



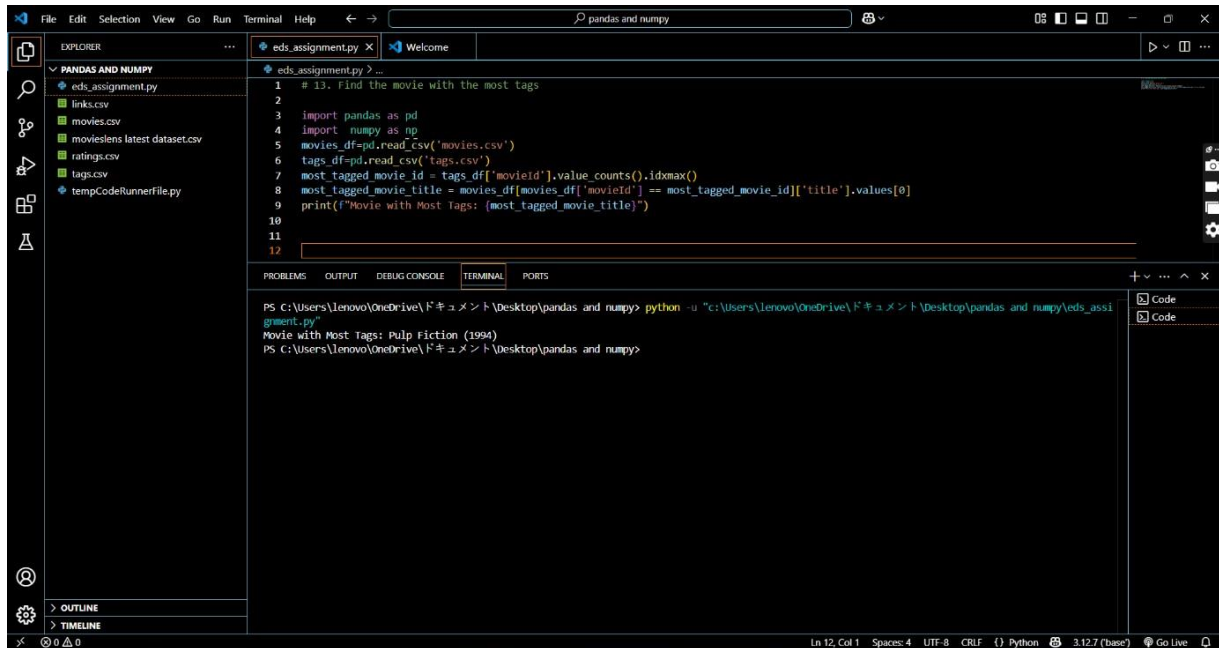
The screenshot shows a VS Code editor window with a file explorer on the left and a code editor in the center. The file explorer shows the same project as the previous screenshot. The code editor shows the following Python code:

```
1 # 12. Top 5 genres with most movies
2
3 import pandas as pd
4 import numpy as np
5 ratings_df = pd.read_csv('ratings.csv')
6 genre_df = pd.read_csv('movies.csv')
7 top5_genres = genre_df['genres'].value_counts().head(5)
8 print(top5_genres)
9
10
11
12
```

The terminal output shows the command to run the script and the result:

```
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"
genres
Drama      1053
Comedy      946
Comedy [Drama]  435
Comedy [Romance]  363
Drama [Romance]  349
Name: count, dtype: int64
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```


13. find the movie with the most tags



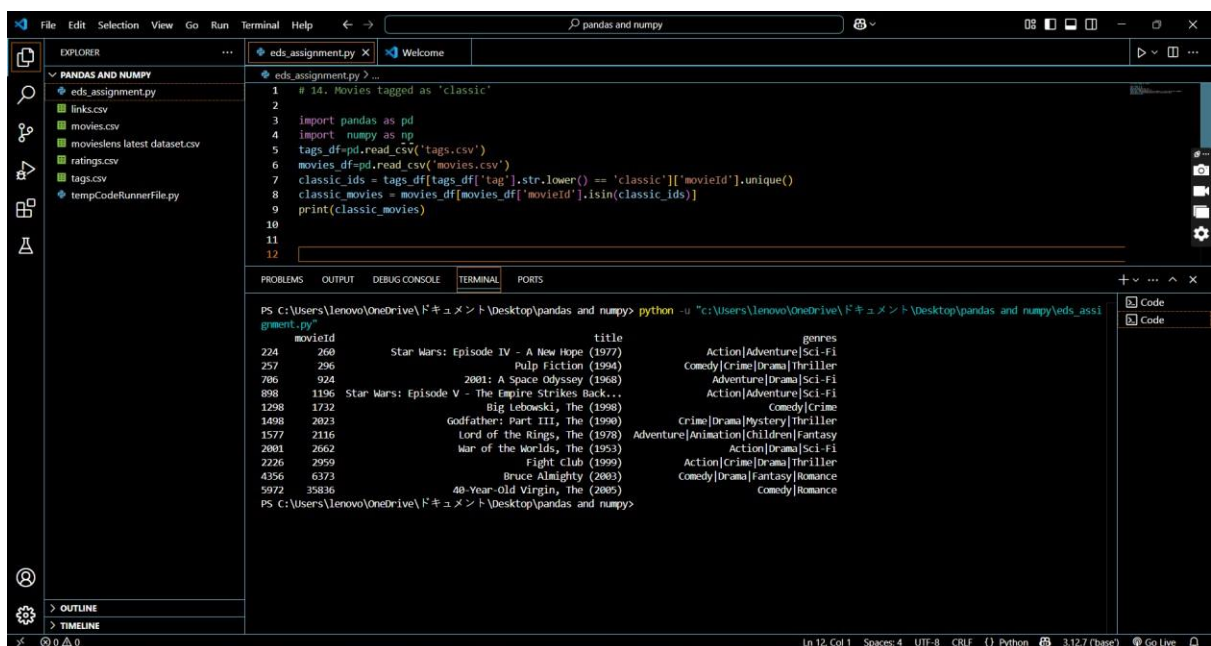
The screenshot shows a VS Code editor with a file explorer on the left containing files like links.csv, movies.csv, ratings.csv, tags.csv, and tempCodeRunnerFile.py. The main editor displays a Python script named eds_assignment.py. The script reads movies.csv and tags.csv, then finds the movie with the most tags by using value counts on the movie_id column of the tags dataframe. The terminal output shows the script executed successfully, identifying 'Pulp Fiction (1994)' as the movie with the most tags.

```
1 # 13. Find the movie with the most tags
2
3 import pandas as pd
4 import numpy as np
5 movies_df=pd.read_csv('movies.csv')
6 tags_df=pd.read_csv('tags.csv')
7 most_tagged_movie_id = tags_df['movieid'].value counts().idxmax()
8 most_tagged_movie_title = movies_df[movies_df['movieid'] == most_tagged_movie_id]['title'].values[0]
9 print(f'Movie with Most Tags: {most_tagged_movie_title}')
10
11
12
```

Terminal Output:

```
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"
Movie with Most Tags: Pulp Fiction (1994)
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```

14. Movies tagged as 'classic'



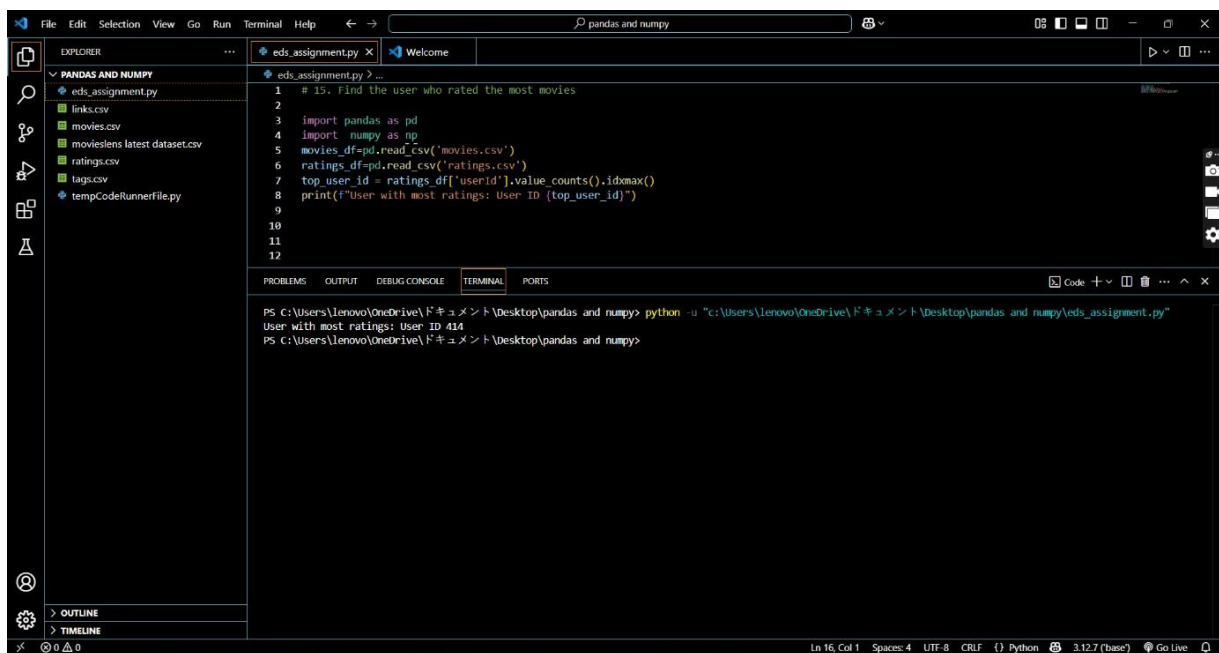
The screenshot shows a VS Code editor with the same file explorer as before. The main editor displays a Python script named eds_assignment.py. The script reads tags.csv and movies.csv, filters for movies tagged as 'classic', and prints a table of these movies along with their genres. The terminal output shows the script executed successfully, displaying a list of classic movies and their genres.

```
1 # 14. Movies tagged as 'classic'
2
3 import pandas as pd
4 import numpy as np
5 tags_df=pd.read_csv('tags.csv')
6 movies_df=pd.read_csv('movies.csv')
7 classic_ids = tags_df[tags_df['tag'].str.lower() == 'classic']['movieid'].unique()
8 classic_movies = movies_df[movies_df['movieid'].isin(classic_ids)]
9 print(classic_movies)
10
11
12
```

Terminal Output:

```
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"
movieid      title                                     genres
224      260      Star Wars: Episode IV - A New Hope      Action|Adventure|Sci-Fi
257      296      Pulp Fiction (1994)                  Comedy|Crime|Drama|Thriller
786      924      2001: A Space Odyssey (1968)                  Adventure|Drama|Sci-Fi
898      1196      Star Wars: Episode V - The Empire Strikes Back...      Action|Adventure|Sci-Fi
1298      1732      Big Lebowski, The (1998)                  Comedy|Crime
1498      2023      Godfather: Part III, The (1990)            Crime|Drama|Mystery|Thriller
1577      2116      Lord of the Rings, The (1978)              Adventure|Animation|Children|Fantasy
2081      2662      War of the Worlds, The (1993)              Action|Drama|Sci-Fi
2226      2959      Fight Club (1999)                        Action|Crime|Drama|Thriller
4356      6373      Bruce Almighty (2003)                     Comedy|Drama|Fantasy|Romance
5972      35836      40-Year-Old Virgin, The (2005)              Comedy|Romance
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```


15. Which user has rated the most movies?

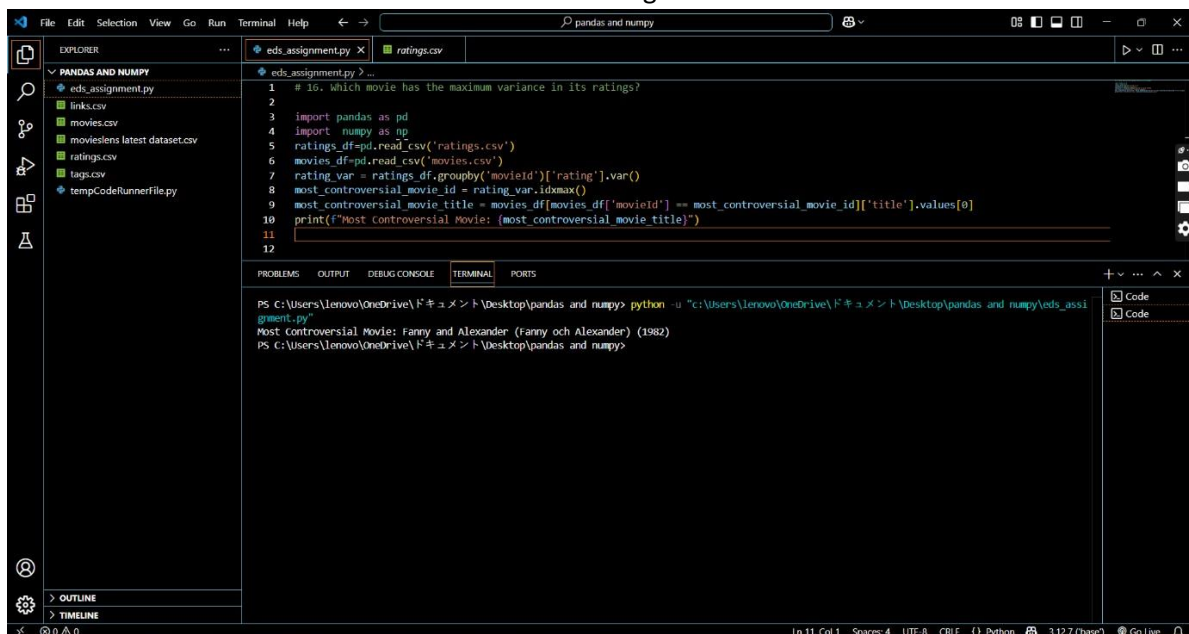


The screenshot shows a VS Code editor with a file explorer on the left containing files like `links.csv`, `movies.csv`, `ratings.csv`, and `tags.csv`. The main editor displays a Python script named `eds_assignment.py` with the following code:

```
1 # 15. Find the user who rated the most movies
2
3 import pandas as pd
4 import numpy as np
5 movies_df = pd.read_csv('movies.csv')
6 ratings_df = pd.read_csv('ratings.csv')
7 top_user_id = ratings_df['userId'].value_counts().idxmax()
8 print(f'User with most ratings: User ID {top_user_id}')
9
10
11
12
```

The terminal at the bottom shows the command `python -u "c:\Users\lenovo\OneDrive\Desktop\pandas and numpy\eds_assignment.py"` and the output `User with most ratings: User ID 414`.

16. which movie has the maximum variance in its rating?

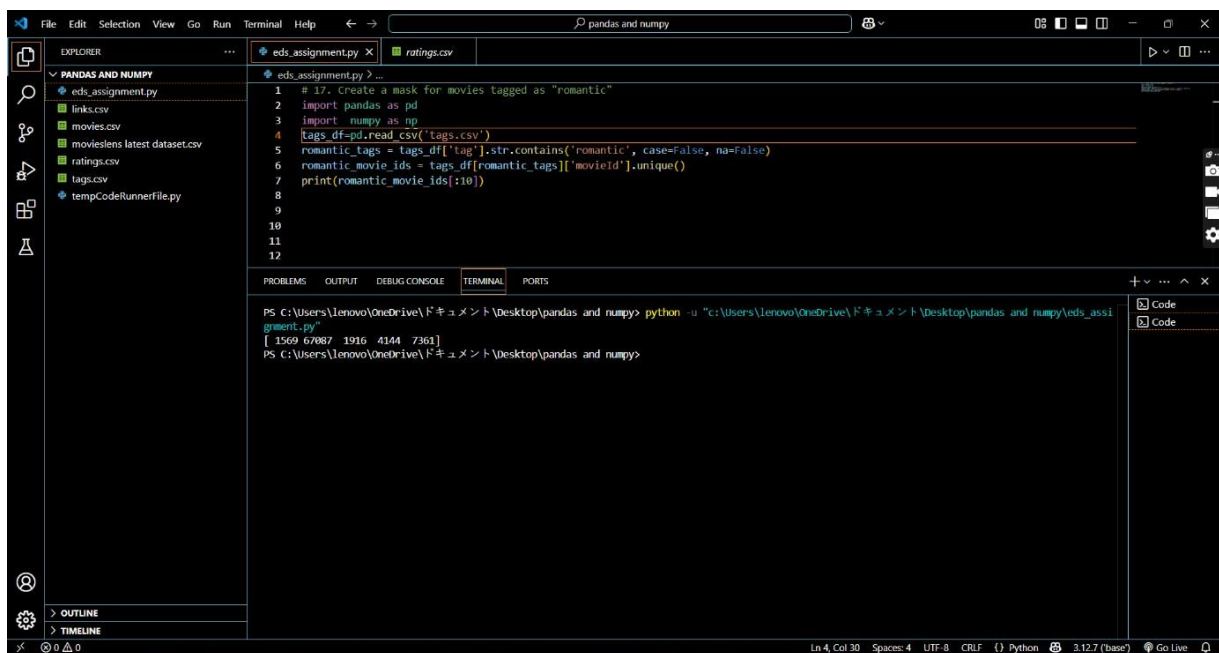


The screenshot shows a VS Code editor with a file explorer on the left. The main editor displays a Python script named `eds_assignment.py` with the following code:

```
1 # 16. Which movie has the maximum variance in its ratings?
2
3 import pandas as pd
4 import numpy as np
5 ratings_df = pd.read_csv('ratings.csv')
6 movies_df = pd.read_csv('movies.csv')
7 rating_var = ratings_df.groupby('movieId')['rating'].var()
8 most_controversial_movie_id = rating_var.idxmax()
9 most_controversial_movie_title = movies_df[movies_df['movieId'] == most_controversial_movie_id]['title'].values[0]
10 print(f'Most controversial Movie: {most_controversial_movie_title}')
11
12
```

The terminal at the bottom shows the command `python -u "c:\Users\lenovo\OneDrive\Desktop\pandas and numpy\eds_assignment.py"` and the output `Most controversial Movie: Fanny och Alexander (fanny och Alexander) (1982)`.

17. create a mask for movies tagged as "romantic"



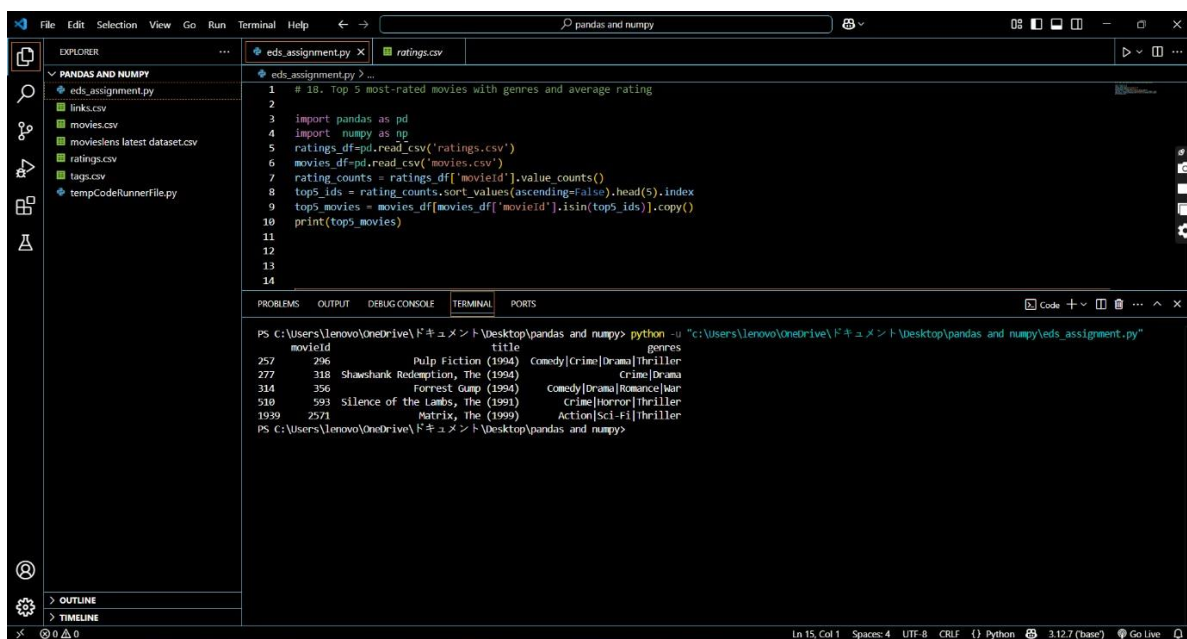
The screenshot shows a VS Code editor with a file explorer on the left containing files like `links.csv`, `movies.csv`, `movielens latest dataset.csv`, `ratings.csv`, and `tags.csv`. The main editor displays a Python script in `eds_assignment.py` with the following code:

```
1 # 17. Create a mask for movies tagged as "romantic"
2 import pandas as pd
3 import numpy as np
4 tags_df = pd.read_csv('tags.csv')
5 romantic_tags = tags_df[tags_df['tag'].str.contains('romantic', case=False, na=False)]
6 romantic_movie_ids = tags_df[romantic_tags]['movieid'].unique()
7 print(romantic_movie_ids[:10])
8
9
10
11
12
```

The terminal at the bottom shows the command `python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"` and its output:

```
[ 1569 67087 1916 4144 7361]
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```

18. Top 5 most-rated movies with genres and average rating



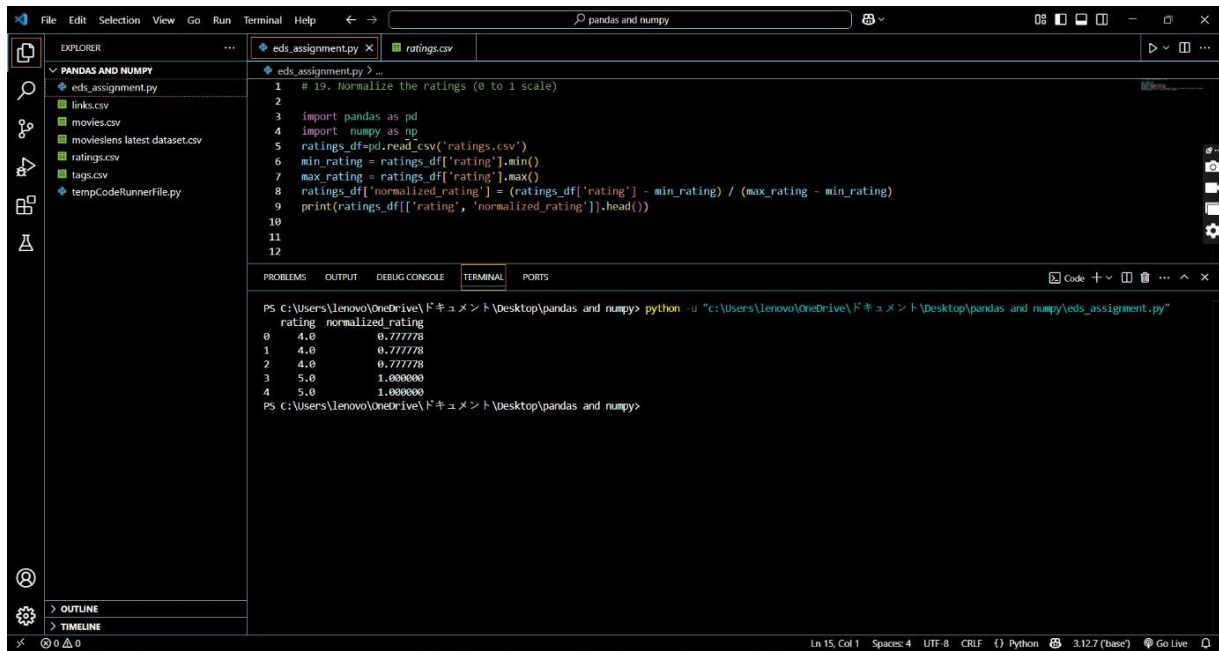
The screenshot shows a VS Code editor with a file explorer on the left containing files like `links.csv`, `movies.csv`, `movielens latest dataset.csv`, `ratings.csv`, and `tags.csv`. The main editor displays a Python script in `eds_assignment.py` with the following code:

```
1 # 18. Top 5 most-rated movies with genres and average rating
2
3 import pandas as pd
4 import numpy as np
5 ratings_df = pd.read_csv('ratings.csv')
6 movies_df = pd.read_csv('movies.csv')
7 rating_counts = ratings_df['movieid'].value_counts()
8 top5_ids = rating_counts.sort_values(ascending=False).head(5).index
9 top5_movies = movies_df[movies_df['movieid'].isin(top5_ids)].copy()
10 print(top5_movies)
11
12
13
14
```

The terminal at the bottom shows the command `python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"` and its output:

```
movieid title genres
257 296 Pulp Fiction (1994) Comedy|Crime|Drama|Thriller
277 318 Shawshank Redemption, The (1994) Crime|Drama
314 356 Forrest Gump (1994) Comedy|Drama|Romance|War
510 593 Silence of the Lambs, The (1991) Crime|Horror|Thriller
1939 2571 Matrix, The (1999) Action|Sci-Fi|Thriller
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```

19. Normalize the ratings (0 to 1 scale)



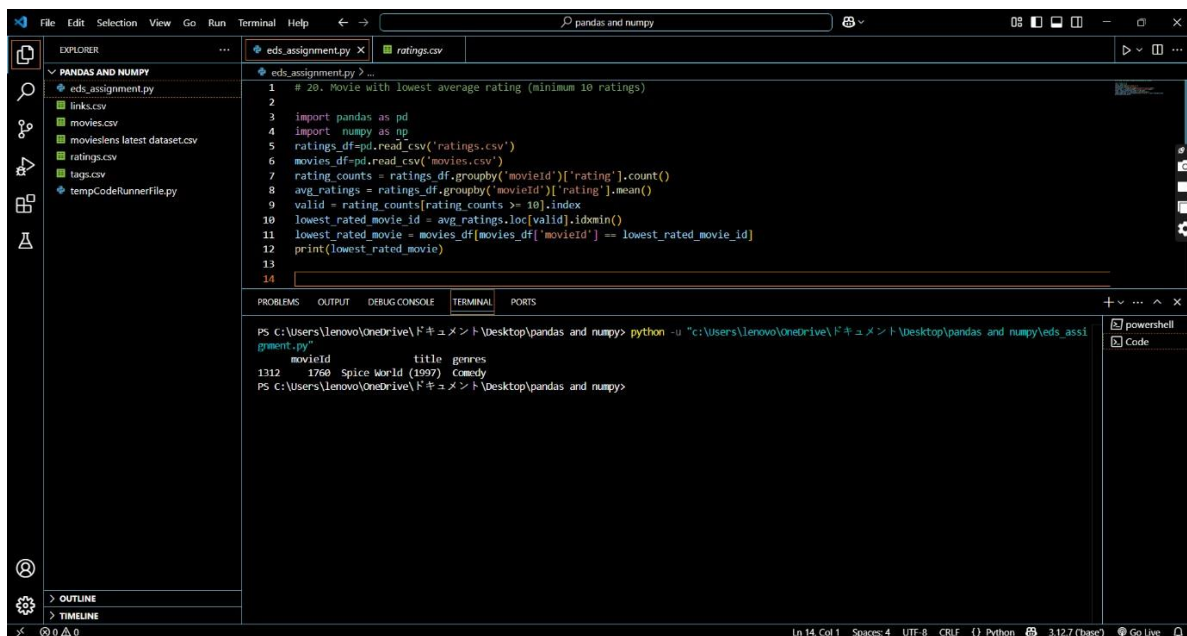
The screenshot shows a VS Code editor with a file explorer on the left containing files like `links.csv`, `movies.csv`, `movieslens latest dataset.csv`, `ratings.csv`, `tags.csv`, and `tempCodeRunnerFile.py`. The main editor displays `eds_assignment.py` with the following code:

```
1 # 19. Normalize the ratings (0 to 1 scale)
2
3 import pandas as pd
4 import numpy as np
5 ratings_df = pd.read_csv('ratings.csv')
6 min_rating = ratings_df['rating'].min()
7 max_rating = ratings_df['rating'].max()
8 ratings_df['normalized_rating'] = (ratings_df['rating'] - min_rating) / (max_rating - min_rating)
9 print(ratings_df[['rating', 'normalized_rating']].head())
10
11
12
```

The terminal at the bottom shows the execution of the script, displaying the first five rows of the normalized ratings:

```
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"
rating normalized_rating
0 4.0 0.777778
1 4.0 0.777778
2 4.0 0.777778
3 5.0 1.000000
4 5.0 1.000000
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```

20. Movie with lowest average rating (minimum 10 ratings)



The screenshot shows the same VS Code editor with `eds_assignment.py` updated with the following code:

```
1 # 20. Movie with lowest average rating (minimum 10 ratings)
2
3 import pandas as pd
4 import numpy as np
5 ratings_df = pd.read_csv('ratings.csv')
6 movies_df = pd.read_csv('movies.csv')
7 rating_counts = ratings_df.groupby('movieId')['rating'].count()
8 avg_ratings = ratings_df.groupby('movieId')['rating'].mean()
9 valid = rating_counts[rating_counts >= 10].index
10 lowest_rated_movie_id = avg_ratings.loc[valid].idxmin()
11 lowest_rated_movie = movies_df[movies_df['movieId'] == lowest_rated_movie_id]
12 print(lowest_rated_movie)
13
14
```

The terminal shows the output of the script, identifying the movie with the lowest average rating:

```
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy> python -u "c:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy\eds_assignment.py"
movieId title genres
1312 1760 Spice World (1997) Comedy
PS C:\Users\lenovo\OneDrive\ドキュメント\Desktop\pandas and numpy>
```